

# Data Wrangling Lab

Estimated time needed: **45 to 60** minutes

In this assignment you will be performing data wrangling.

## Objectives

In this lab you will perform the following:

- Identify duplicate values in the dataset.
- Remove duplicate values from the dataset.
- Identify missing values in the dataset.
- Impute the missing values in the dataset.
- Normalize data in the dataset.

---

## Hands on Lab

Import pandas module.

```
In [1]: import pandas as pd
```

Load the dataset into a dataframe.

```
In [2]: df = pd.read_csv("https://cf-courses-data.s3.us.cloud-object-storage.appdomain.cloud/IBM-DA0321EN-SkillsNetwork/LargeData/m1_")
```

```
In [3]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 11552 entries, 0 to 11551

Data columns (total 85 columns):

#	Column	Non-Null Count	Dtype
0	Respondent	11552 non-null	int64
1	MainBranch	11552 non-null	object
2	Hobbyist	11552 non-null	object
3	OpenSourcer	11552 non-null	object
4	OpenSource	11471 non-null	object
5	Employment	11552 non-null	object
6	Country	11552 non-null	object
7	Student	11499 non-null	object
8	EdLevel	11436 non-null	object
9	UndergradMajor	10812 non-null	object
10	EduOther	11388 non-null	object
11	OrgSize	11454 non-null	object
12	DevType	11485 non-null	object
13	YearsCode	11543 non-null	object
14	Age1stCode	11539 non-null	object
15	YearsCodePro	11536 non-null	object
16	CareerSat	11552 non-null	object
17	JobSat	11551 non-null	object
18	MgrIdiot	11054 non-null	object
19	MgrMoney	11050 non-null	object
20	MgrWant	11054 non-null	object
21	JobSeek	11552 non-null	object
22	LastHireDate	11552 non-null	object
23	LastInt	11129 non-null	object
24	FizzBuzz	11515 non-null	object
25	JobFactors	11549 non-null	object
26	ResumeUpdate	11511 non-null	object
27	CurrencySymbol	11552 non-null	object
28	CurrencyDesc	11552 non-null	object
29	CompTotal	10737 non-null	float64
30	CompFreq	11346 non-null	object
31	ConvertedComp	10730 non-null	float64
32	WorkWeekHrs	11427 non-null	float64
33	WorkPlan	11429 non-null	object
34	WorkChallenge	11384 non-null	object
35	WorkRemote	11544 non-null	object
36	WorkLoc	11520 non-null	object
37	ImpSyn	11547 non-null	object
38	CodeRev	11551 non-null	object
39	CodeRevHrs	9083 non-null	float64
40	UnitTests	11523 non-null	object
41	PurchaseHow	11354 non-null	object
42	PurchaseWhat	11514 non-null	object
43	LanguageWorkedWith	11541 non-null	object
44	LanguageDesireNextYear	11415 non-null	object
45	DatabaseWorkedWith	11096 non-null	object
46	DatabaseDesireNextYear	10497 non-null	object
47	PlatformWorkedWith	11130 non-null	object
48	PlatformDesireNextYear	10991 non-null	object
49	WebFrameWorkedWith	10139 non-null	object
50	WebFrameDesireNextYear	9918 non-null	object
51	MiscTechWorkedWith	9343 non-null	object
52	MiscTechDesireNextYear	10078 non-null	object
53	DevEnviron	11523 non-null	object
54	OpSys	11518 non-null	object
55	Containers	11470 non-null	object
56	BlockchainOrg	9198 non-null	object
57	BlockchainIs	8915 non-null	object
58	BetterLife	11452 non-null	object
59	ITperson	11517 non-null	object
60	Off0n	11514 non-null	object
61	SocialMedia	11251 non-null	object
62	Extraversion	11532 non-null	object
63	ScreenName	11039 non-null	object
64	S0Visit1st	11227 non-null	object
65	S0VisitFreq	11547 non-null	object
66	S0VisitTo	11551 non-null	object
67	S0FindAnswer	11549 non-null	object
68	S0TimeSaved	11501 non-null	object
69	S0HowMuchTime	9616 non-null	object
70	S0Account	11551 non-null	object
71	S0PartFreq	10404 non-null	object
72	S0Jobs	11546 non-null	object
73	EntTeams	11547 non-null	object
74	S0Comm	11552 non-null	object
75	WelcomeChange	11463 non-null	object
76	S0NewContent	9557 non-null	object
77	Age	11255 non-null	float64
78	Gender	11477 non-null	object
79	Trans	11429 non-null	object
80	Sexuality	11005 non-null	object

```
81 Ethnicity          10869 non-null object
82 Dependents         11408 non-null object
83 SurveyLength       11533 non-null object
84 SurveyEase         11538 non-null object
dtypes: float64(5), int64(1), object(79)
memory usage: 7.5+ MB
```

## Finding duplicates

In this section you will identify duplicate values in the dataset.

Find how many duplicate rows exist in the dataframe.

```
In [4]: duplicate_rows = df[df.duplicated()]
len(duplicate_rows)
```

Out[4]: 154

```
In [5]: #How many duplicate values are there in the column Respondent?
len(df['Respondent'])-len(df['Respondent'].drop_duplicates())
```

Out[5]: 154

## Removing duplicates

Remove the duplicate rows from the dataframe.

```
In [6]: df.drop_duplicates(keep='first', inplace=True)
```

Verify if duplicates were actually dropped.

```
In [7]: duplicate_rows = df[df.duplicated()]
len(duplicate_rows)
```

Out[7]: 0

```
In [8]: #After removing the duplicate rows, how many rows are there in the dataset?
len(df)
```

Out[8]: 11398

```
In [9]: #After removing the duplicate rows, how many unique rows are there in the column Respondent?
len(df['Respondent'])
```

Out[9]: 11398

## Finding Missing values

Find the missing values for all columns.

```
In [10]: missing_data = df.isnull() #True means the value is missing
missing_data.head(5)
```

```
Out[10]:
```

	Respondent	MainBranch	Hobbyist	OpenSourcer	OpenSource	Employment	Country	Student	EdLevel	UndergradMajor	...	WelcomeCha
0	False	False	False	False	False	False	False	False	False	False	...	F
1	False	False	False	False	False	False	False	False	False	False	...	F
2	False	False	False	False	False	False	False	False	False	False	...	F
3	False	False	False	False	False	False	False	False	False	True	...	F
4	False	False	False	False	False	False	False	False	False	False	...	F

5 rows × 85 columns

Find out how many rows are missing in the column 'EdLevel' and 'Country'

```
In [11]: missing_data['EdLevel'].value_counts()
```

```
Out[11]: False    11286
          True      112
          Name: EdLevel, dtype: int64
```

```
In [12]: missing_data['Country'].value_counts()
```

```
Out[12]: False      11398
Name: Country, dtype: int64
```

## Imputing missing values

Find the value counts for the column WorkLoc.

```
In [13]: df['WorkLoc'].value_counts()
```

```
Out[13]: Office      6806
Home      3589
Other place, such as a coworking space or cafe    971
Name: WorkLoc, dtype: int64
```

Identify the value that is most frequent (majority) in the WorkLoc column.

```
In [14]: df['WorkLoc'].value_counts().idxmax()
```

```
Out[14]: 'Office'
```

```
In [15]: #What is the majority category under the column Employment?
df['Employment'].value_counts().idxmax()
```

```
Out[15]: 'Employed full-time'
```

```
In [16]: #Under the column " UndergradMajor", which category has the minimum number of rows?
df['UndergradMajor'].value_counts().idxmin()
```

```
Out[16]: 'A health science (ex. nursing, pharmacy, radiology)'
```

Impute (replace) all the empty rows in the column WorkLoc with the value that you have identified as majority.

```
In [17]: import numpy as np
df['WorkLoc'].replace(np.nan, 'Office', inplace=True)
```

After imputation there should ideally not be any empty rows in the WorkLoc column.

Verify if imputing was successful.

```
In [18]: missing_data = df.isnull()
len(df['WorkLoc']) == len(missing_data[missing_data['WorkLoc'] == False])
```

```
Out[18]: True
```

## Normalizing data

There are two columns in the dataset that talk about compensation.

One is "CompFreq". This column shows how often a developer is paid (Yearly, Monthly, Weekly).

The other is "CompTotal". This column talks about how much the developer is paid per Year, Month, or Week depending upon his/her "CompFreq".

This makes it difficult to compare the total compensation of the developers.

In this section you will create a new column called 'NormalizedAnnualCompensation' which contains the 'Annual Compensation' irrespective of the 'CompFreq'.

Once this column is ready, it makes comparison of salaries easy.

---

List out the various categories in the column 'CompFreq'

```
In [19]: df['CompFreq'].value_counts()
```

```
Out[19]: Yearly      6073
Monthly    4788
Weekly     331
Name: CompFreq, dtype: int64
```

```
In [20]: #delete the rows in which CompFreq or CompTotal is missing
df.dropna(subset=['CompFreq', 'CompTotal'], axis=0, inplace=True)
```

Create a new column named 'NormalizedAnnualCompensation'. Use the hint given below if needed.

Double click to see the **Hint**.

```
In [21]: #create a new empty column named 'NormalizedAnnualCompensation'
df['NormalizedAnnualCompensation'] = ""

#If CompFreq is Yearly, use the existing value in CompTotal
#If CompFreq is Monthly, multiply the value in CompTotal by 12
#If CompFreq is Weekly, multiply the value in CompTotal by 52 (weeks in an year)
conditions = [(df['CompFreq'] == 'Yearly'), (df['CompFreq'] == 'Monthly'), (df['CompFreq'] == 'Weekly')]
choices = [df['CompTotal'], df['CompTotal']*12, df['CompTotal']*52]
df['NormalizedAnnualCompensation'] = np.select(conditions, choices, default=0)

df[['CompFreq', 'CompTotal', 'NormalizedAnnualCompensation']].head(15)
```

Out [21]:

	CompFreq	CompTotal	NormalizedAnnualCompensation
0	Yearly	61000.0	61000.0
1	Yearly	138000.0	138000.0
2	Yearly	90000.0	90000.0
3	Monthly	29000.0	348000.0
4	Yearly	90000.0	90000.0
5	Monthly	9500.0	114000.0
6	Monthly	3000.0	36000.0
7	Yearly	103000.0	103000.0
8	Yearly	69000.0	69000.0
9	Monthly	8000.0	96000.0
10	Monthly	7000.0	84000.0
11	Yearly	114000.0	114000.0
12	Weekly	2000.0	104000.0
13	Weekly	22000.0	1144000.0
14	Monthly	96000.0	1152000.0

```
In [22]: #What is the median NormalizedAnnualCompensation?
df['NormalizedAnnualCompensation'].median()
```

Out [22]: 100000.0

## Authors

Ramesh Sannareddy

## Other Contributors

Rav Ahuja

## Change Log

Date (YYYY-MM-DD)	Version	Changed By	Change Description
2020-10-17	0.1	Ramesh Sannareddy	Created initial version of the lab

Copyright © 2020 IBM Corporation. This notebook and its source code are released under the terms of the [MIT License](#).